

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split, KFold, cross_val_score, cross_validate
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score, make_scorer
```

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

✓ Step 1: Input the processed data

```
data = pd.read_csv('/content/drive/MyDrive/Data testing with IDXExchange/data_filtered-6nd-log_ClosePrice.csv')
```

```
data.info()
```

↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 79543 entries, 0 to 79542
Data columns (total 17 columns):

#	Column	Non-Null Count	Dtype
0	AttachedGarageYN	79543 non-null	float64
1	BathroomsTotalInteger	79543 non-null	float64
2	BedroomsTotal	79543 non-null	float64
3	FireplaceYN	79543 non-null	float64
4	GarageSpaces	79543 non-null	float64
5	Latitude	79543 non-null	float64
6	LivingArea	79543 non-null	float64
7	Longitude	79543 non-null	float64
8	LotSizeSquareFeet	79543 non-null	float64
9	NewConstructionYN	79543 non-null	float64
10	ParkingTotal	79543 non-null	float64
11	PoolPrivateYN	79543 non-null	float64
12	Stories	79543 non-null	float64
13	ViewYN	79543 non-null	float64
14	ListingContractDate_month	79543 non-null	float64
15	Age	79543 non-null	float64
16	log_ClosePrice	79543 non-null	float64

dtypes: float64(17)
memory usage: 10.3 MB

✓ Train-Test Split

Features/Target:

- x: All features except log_ClosePrice
- y: Transformed target (log_ClosePrice)

Split:

80% train / 20% test

- Random state fixed (42) for reproducibility

Purpose:

Creates evaluation framework for model development

```
X = data.drop(['log_ClosePrice'], axis=1)
y = data['log_ClosePrice']
```

```
# Step 2: Random train-test split (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

✓ Modeling Pipeline

Components:

1. StandardScaler : Normalizes all features
2. RandomForestRegressor :
 - 30 trees (n_estimators=30)
 - Fixed random state (42) for reproducibility

Design:

Single pipeline ensures:

- Consistent preprocessing for train/test
- Avoids data leakage
- Simplified model deployment

```
# Step 3: Define pipeline
pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('rf', RandomForestRegressor(n_estimators=30, random_state=42))
])
```

✓ Model Validation

Method:

5-fold cross-validation with:

- Shuffled splits (shuffle=True)
- Fixed random state (42)

Metrics Reported:

- **RMSE:**
{mean:.2f} ± {std:.2f}
- **MAE:**
{mean:.2f}
- **R²:**
{mean:.4f}

Purpose:

Robust performance estimation before final test evaluation

```
# Step 4: Cross-validation on training set
kf = KFold(n_splits=5, shuffle=True, random_state=42)

mse_scores = cross_val_score(pipeline, X_train, y_train, cv=kf, scoring='neg_mean_squared_error')
rmse_scores = np.sqrt(-mse_scores)

print("Cross-validated RMSE scores on Training Set:", rmse_scores)
print(f"Average RMSE: {rmse_scores.mean():.2f}")
print(f"Standard Deviation of RMSE: {rmse_scores.std():.2f}")

🔗 Cross-validated RMSE scores on Training Set: [0.21286903 0.18018384 0.17441161 0.18330647 0.17211637]
Average RMSE: 0.18
Standard Deviation of RMSE: 0.01

# Optional: Other metrics
scoring = {
    'mae': make_scorer(mean_absolute_error),
    'mse': make_scorer(mean_squared_error),
    'r2': make_scorer(r2_score)
}

results = cross_validate(pipeline, X_train, y_train, cv=kf, scoring=scoring)

print(f"MAE scores: {results['test_mae']}")
print(f"Average MAE: {np.mean(results['test_mae']):.2f}")

print(f"RMSE scores: {np.sqrt(results['test_mse'])}")
print(f"Average RMSE: {np.mean(np.sqrt(results['test_mse'])):.2f}")
```

```
print(f"R² scores: {results['test_r2']}")
print(f"Average R²: {np.mean(results['test_r2']):.4f}")

MAE scores: [0.11683174 0.1169057 0.11387609 0.11552286 0.11511504]
Average MAE: 0.12
RMSE scores: [0.21286903 0.18018384 0.17441161 0.18330647 0.17211637]
Average RMSE: 0.18
R² scores: [0.84315796 0.88446971 0.88799685 0.88116587 0.89167799]
Average R²: 0.8777
```

Final Model Evaluation

Process:

1. Trained pipeline on full training set
2. Predicted on held-out test set

Test Metrics:

- **MAE:** {value:.4f}
- **MSE:** {value:.4f}
- **RMSE:** {value:.4f}
- **R²:** {value:.4f}

Purpose:

Unbiased performance estimate of production-ready model

```
# Step 5: Final model training and evaluation on test set
pipeline.fit(X_train, y_train)
y_pred_test = pipeline.predict(X_test)

mae_test = mean_absolute_error(y_test, y_pred_test)
mse_test = mean_squared_error(y_test, y_pred_test)
rmse_test = np.sqrt(mse_test)
r2_test = r2_score(y_test, y_pred_test)

print("\n--- Final Evaluation on Test Set ---")
print(f"MAE: {mae_test:.4f}")
print(f"MSE: {mse_test:.4f}")
print(f"RMSE: {rmse_test:.4f}")
print(f"R² Score: {r2_test:.4f}")
```

```
--- Final Evaluation on Test Set ---
MAE: 0.1131
MSE: 0.0295
RMSE: 0.1718
R² Score: 0.8907
```

Prediction Visualization

Plot Components:

- Points: Actual vs Predicted values (log_ClosePrice)
- Red line: Perfect prediction reference

Key Insights:

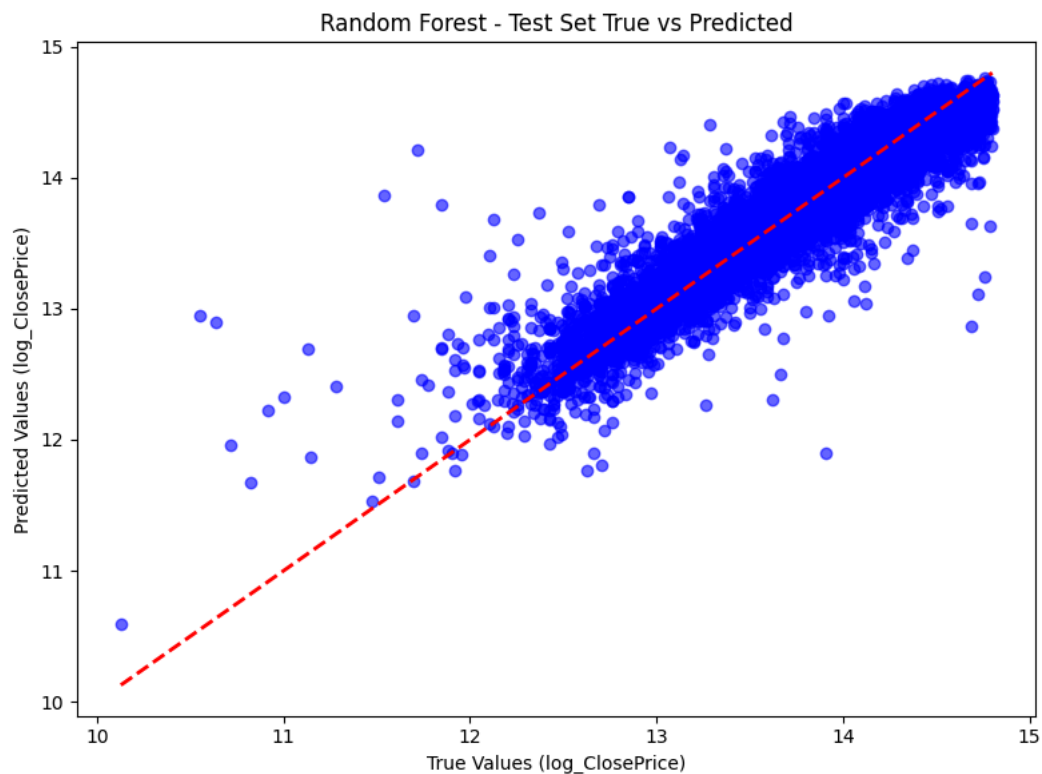
- Vertical spread shows prediction error magnitude
- Cloud tightness indicates model accuracy
- Outliers visible as distant points

Purpose:

Visual diagnostic beyond numeric metrics

```
# Step 6: Plot True vs Predicted on Test Set
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred_test, color='blue', alpha=0.6)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linestyle='--', lw=2)
plt.xlabel('True Values (log_ClosePrice)')
plt.ylabel('Predicted Values (log_ClosePrice)')
plt.title('Random Forest - Test Set True vs Predicted')
```

```
plt.tight_layout()
plt.show()
```



✓ USD Price Prediction Plot

Transformation:

Converted predictions from log-scale back to original USD prices

Key Features:

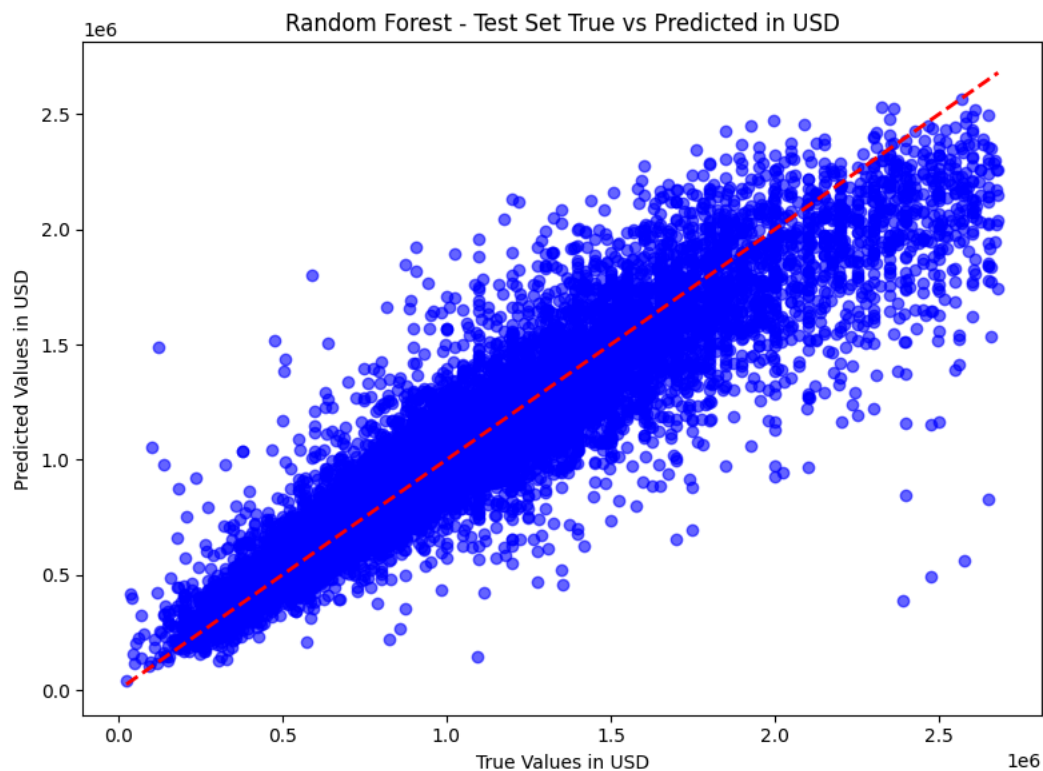
- Points show actual vs predicted home values
- Red line represents perfect predictions
- Axes now in interpretable USD

Business Value:

Enables direct evaluation of dollar-value accuracy

```
# Convert from log scale back to original price scale
y_test_actual = np.exp(y_test)
y_pred_actual = np.exp(y_pred_test)

# Step 6: Plot True vs Predicted on Test Set
plt.figure(figsize=(8, 6))
plt.scatter(y_test_actual, y_pred_actual, color='blue', alpha=0.6)
plt.plot([min(y_test_actual), max(y_test_actual)], [min(y_test_actual), max(y_test_actual)], color='red', linestyle='--', lw=2)
plt.xlabel('True Values in USD')
plt.ylabel('Predicted Values in USD')
plt.title('Random Forest - Test Set True vs Predicted in USD')
plt.tight_layout()
plt.show()
```



Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.